



Indian Institute of Technology Bhilai  
Department of Computer Science  
CSL301: Operating Systems  
Scope: xv6  
Difficulty Level: Moderate

Take Home  
Assignment - 2  
August 28, 2025

## Instructions -

- Follow the steps given in the question and update all required files
- Create a Word (.doc/docx) file with the following (Use LibreOffice in Ubuntu) in step-by-step format as mentioned in the question. Provide only the code that you have added and the file name under which the update has been made.
- Save your final shell source code with the following naming convention: `<rollnumber>_sh.c`
- Do not include full files, just the specific changes made. Save this file using: `<rollnumber>_part.pdf`
- Submit these files as follows: Place all the files above in a single folder and compress it. Name the compressed archive as `HA2_<rollnumber>_part.zip`

**Note:** Do not include explanations or code outside the specified text file and annotated source file. Each screenshot should clearly show your QEMU terminal running and reflecting the update for each task.

## Objective

- Implement a new system call `pinfo(int pid, struct proc_info *info)` in xv6.
- Retrieve process information including PID, name, state, and memory size.
- Practice modifying both kernel and user-space code.
- Test your implementation using multiple processes generated by a helper program.

## Background and Hints

Each process in xv6 is represented by a `struct proc` in `kernel/proc.h`. Relevant fields include:

- `pid` – the process ID
- `name` – name of the process (string)
- `state` – process state (`UNUSED`, `USED`, `SLEEPING`, `RUNNABLE`, `RUNNING`, `ZOMBIE`)
- `sz` – memory size of the process

## Safe Locking in xv6

In xv6, the process table (`ptable.proc[]`) is a shared resource accessed by multiple processes. To prevent race conditions, **safe locking** must be used with `ptable.lock`.

**Note:** Safe locking is a core concept in concurrency and will be taught in your concurrency class. This assignment gives a practical example of using locks in the kernel.

### 1. Files that require `spinlock.h`:

- `proc.c` – Required to acquire/release `ptable.lock` in `get_proc_info()`.
- `sysproc.c` – Required if accessing `ptable` or calling functions that acquire locks.
- `proc.h` – Optional; include only if declaring functions or structures that reference locks.
- Other kernel files accessing `ptable` (e.g., `fork.c`, `trap.c`) – include `spinlock.h` if you acquire or release a lock.
- User programs (`pinfo.c`, `testproc.c`) do NOT include `spinlock.h`.

## 2. How to implement safe locking:

1. Include the spinlock header in kernel files:

```
#include "spinlock.h"
#include "proc.h"
```

2. Acquire the lock before accessing the process table:

```
acquire(&ptable.lock);
```

3. Access or modify the shared data (e.g., read PID, name, state, memory size).

4. Release the lock after finishing, even if returning due to an error:

```
release(&ptable.lock);
```

### Example in `get_proc_info`:

```
int
get_proc_info_kernel(int pid, struct proc_info *info)
{
    struct proc *p;

    for(p = proc; p < &proc[NPROC]; p++){
        acquire(&p->lock);
        if(p->pid == pid){
            info->pid = p->pid;
            safestrcpy(info->name, p->name, sizeof(info->name));

            // Convert enum state to string
            char *st = "UNKNOWN";
            switch(p->state){
                case UNUSED:  st = "UNUSED";  break;
                case USED:    st = "USED";    break;
                case SLEEPING: st = "SLEEPING"; break;
                case RUNNABLE: st = "RUNNABLE"; break;
                case RUNNING:  st = "RUNNING"; break;
                case ZOMBIE:   st = "ZOMBIE";  break;
            }
            safestrcpy(info->state, st, sizeof(info->state));

            info->sz = p->sz;

            release(&p->lock);
            return 0;
        }
        release(&p->lock);
    }
    return -1;
}
```

### Key Points:

- Always acquire the lock before reading/modifying shared data.

- Always release the lock, even on errors.
- Use `safestrcpy` for kernel string copying.
- Only kernel files require `spinlock.h`, not user programs.
- Understanding locks is part of your concurrency course; this assignment gives a hands-on example.

## Kernel-side Modifications:

1. Create `procinfo.h` in both `kernel/` and `user/`. Define `struct proc_info` with fields:
  - `int pid;`
  - `char name[16];`
  - `char state[16];`
  - `uint64 sz;`
2. Implement `get_proc_info(int pid, struct proc_info *info)` in `proc.c`.
3. Add syscall number in `syscall.h`, implement `sys_get_proc_info()` in `sysproc.c`, and wire it in `syscall.c`.

## User Program: `pinfo.c`

The user program should:

- Accept a PID as a command-line argument.
- Call `get_proc_info(pid, &info)`.
- Display: PID, Name, State, and Memory Size.
- Handle invalid PIDs gracefully.

## Testing with `testproc.c`

A helper program `testproc.c` should be created to fork multiple processes for testing.

```
#include "user/user.h"

int main(void) {
    int i;
    int num_children = 5;

    for(i = 0; i < num_children; i++) {
        int pid = fork();
        if(pid < 0) {
            printf("Fork failed\n");
            exit(1);
        }
        if(pid == 0) {
            printf("Child process %d started with PID %d\n", i+1, getpid());
            while(1); // never exits
        }
    }

    // Parent waits for all children to finish
    // for(i = 0; i < num_children; i++) {
    //     int wpid = wait(0);
    //     printf("Parent: child PID %d finished\n", wpid);
    // }

    exit(0);
}
```

## Instructions

1. Place `testproc.c` in `user/`.
2. Add `_testproc` to `UPROGS` in the Makefile.
3. Compile `xv6` using `make qemu` and first run `testproc` in `xv6` and then press enter after that run `pinfo`.
4. Run the program inside the `xv6` shell:

```
$ testproc
$ pinfo 3
$ pinfo 4
$ pinfo 5
```

5. Observe how `pinfo` displays the information for multiple processes.

## Expected Output and Testing Guidelines

```
PID: 2
Name: sh
State: RUNNING
Size: 4096
```

### Testing tips:

- Ensure `testproc` is running so children exist when checking with `pinfo`.
- Check multiple PIDs including invalid ones to confirm proper error handling.
- Verify that `state` strings correctly match the process enum (`RUNNING`, `SLEEPING`, etc.).

## Submission and Evaluation

### Deliverables:

- Modified files: `proc.c`, `proc.h`, `sysproc.c`, `syscall.c`, `syscall.h`, `pinfo.c`.
- Screenshot of `pinfo` output for at least 2 processes.

### Evaluation Criteria:

- Correctness of implementation
- Proper state conversion from enum to string
- Safe locking
- Robust error handling
- Clean, readable code